

Drivable by Construction, Behind on Accuracy: A Kinematic Decoder for Transformer Motion Forecasting on Argoverse 2

Ali Taheri

Code and configurations: <https://github.com/tali7965/Physics-Constrained-Transformer-for-Motion-Forecasting>

Abstract—Most learned motion forecasters regress future coordinates directly, which allows trajectories that no vehicle could drive. We test the alternative of making every prediction drivable by construction. The goal is not the most accurate forecaster, but to build a physically accurate one and measure how it performs. A compact Transformer (1.43M parameters) outputs bounded acceleration and steering, which a differentiable kinematic bicycle model integrates into positions. We compare it with the same model using an unconstrained coordinate decoder, under one protocol on Argoverse 2. The kinematic decoder removes every infeasible trajectory but is less accurate: validation brier-minFDE at $K=6$ is 3.16m against 2.57m. The gap, 0.41–0.73m, holds across two seeds, across 10%, 25% and 100% of a 50k-scene training subset, with and without the map, and at $K=1, 3$ and 6, so the physical prior brings no gain in data efficiency. The coordinate decoder’s infeasibility is step-to-step jitter: 58% of its most probable trajectories exceed 8m/s^2 of lateral acceleration at 10 Hz, and none do at 2 Hz. The kinematic decoder also leaves the drivable area more often (5.2% against 1.5%). We trace the accuracy gap to the integration, which gives the first second’s controls about 60 times the gradient of the last second’s. The off-road excess comes from errors across the lane: an initial heading the model does not correct, and turns it does not complete.

I. INTRODUCTION

Motion forecasting predicts where the agents around an autonomous vehicle will be over the next few seconds. Most current forecasters encode the scene with attention over agents and map elements [1]–[5] and regress future positions directly. Nothing in such a decoder ties consecutive positions to vehicle motion, so it can predict sideways slides, turns too sharp for the vehicle’s speed, or jitter between steps. A long-standing remedy is to predict control inputs instead and integrate them through a vehicle model [6], [7], so that every output is kinematically feasible by construction.

The goal of this project is not to build the most accurate forecaster, nor to show that a kinematic decoder is the best approach. It is to build a physically accurate model and measure, in a controlled setting, what that constraint costs and buys. One Transformer encoder is trained twice on Argoverse 2 [8], and only the output head differs. One head regresses 60 future positions. The other outputs 60 bounded (acceleration, steering) pairs that a differentiable kinematic bicycle model [9], [10] integrates. Three questions guide the study: (1) does the kinematic decoder match the coordinate decoder’s accuracy, (2) how often are unconstrained trajectories infeasible, and

does the kinematic decoder remove them, and (3) does the physical prior help more when training data is limited?

The answers are mostly negative for the constrained head, and the reasons are specific:

- In every setting tested (two seeds, three training-set sizes, with and without the map, and three numbers of modes) the kinematic decoder is behind by 0.41–0.73m brier-minFDE. That is about 40 times the seed-to-seed spread, and the gap does not shrink with less data.
- The coordinate decoder’s infeasibility is almost entirely 10 Hz jitter. Resampled to 2 Hz, its trajectories are as feasible as the ground truth’s or more so.
- Integration skews the gradient towards early controls by a factor of about 60, and the kinematic decoder’s accelerations chatter at their bounds. Its off-road excess comes from an uncorrected initial heading and from incomplete turns, not from the chatter.

II. RELATED WORK

Scene encoders. VectorNet [1] and LaneGCN [11] replaced rasterised scenes with vectorised agents and lanes. Transformer forecasters then attend jointly over agents, time and map [2]–[5]. Ours is a small instance with PointNet-style [12] tokens, close to Wayformer’s early-fusion variant [3].

Multimodal outputs. Forecasters predict several modes with probabilities and train them with a winner-takes-all loss [13], sometimes around fixed anchors [14]. CoverNet [15] classifies over a set of trajectories that can be made feasible in advance.

Kinematic constraints. DKM [6] adds a kinematic layer that turns a network’s acceleration and steering outputs into vehicle trajectories. Trajectron++ [7] integrates predicted controls through agent-specific dynamics. The kinematic bicycle model is standard for planning [9] and stays consistent with vehicle motion while lateral acceleration is moderate [10]. We use it in the same way but isolate its effect: the encoder, losses, data and selection rule are shared, only the head changes, and feasibility and off-road rates are measured next to accuracy.

III. METHOD

A. Task and data

Argoverse 2 motion forecasting [8] gives 5 s of history for all agents and a vector map, and asks for the focal agent’s next 6 s at 10 Hz (60 steps). Each scene is moved to a frame

centred on the focal agent at the last observed step, with the x -axis along its recorded heading. It keeps up to 64 agents and 192 lane segments within 150m, each lane as a 10-point centreline.

Protocol. A seeded permutation of the training split gives a 5,000-scene *dev* set, used only for model selection, and a fixed 50,000-scene training subset. Its first 10% and 25% form nested subsets for the data-efficiency study. The official validation split (24,988 scenes) is used only for reporting. The challenge’s single-agent leaderboard stopped accepting submissions on 31 October 2025, so no test results are given.

B. Encoder

An agent’s history steps (position, velocity, cos and sin of heading) are embedded linearly, a learned timestep embedding is added, a LayerNorm–ReLU–Linear block follows, and the result is max-pooled over valid steps into one token, plus a type embedding. A lane’s 9 centreline segments (start and end points) pass through an MLP and are max-pooled, with lane-type and intersection embeddings added. A 4-layer pre-norm Transformer encoder [16] ($d=128$, 8 heads, FFN 512) runs self-attention over all 256 agent and lane tokens, with empty slots masked. There is no slot positional encoding, so the output does not depend on slot order. K mode queries, each the focal token plus a learned mode embedding, attend to the scene in a 2-layer Transformer decoder. An MLP head maps each mode to a trajectory, and a linear layer gives its logit. The model has 1.43M parameters.

C. Coordinate and kinematic heads

The *coordinate head* is an MLP that outputs 60 positions per mode.

The *kinematic head* (the physics head in the code) has the same MLP but outputs 60 raw pairs (u_k^a, u_k^s) , bounded to an acceleration $a_k = a_{\max} \tanh u_k^a$ and a steering fraction $s_k = \tanh u_k^s$. Curvature is $\kappa_k = s_k \min(\kappa_{\max}, a_{\text{lat}}/\bar{v}_k^2)$: a 5 m turning radius at low speed and at most a_{lat} of lateral acceleration above it. Each step of $\Delta t = 0.1$ s integrates

$$\begin{aligned} v_{k+1} &= \max(v_k + a_k \Delta t, 0), \quad \bar{v}_k = \frac{1}{2}(v_k + v_{k+1}), \\ \Delta s_k &= \bar{v}_k \Delta t, \quad \bar{\psi}_k = \psi_k + \frac{1}{2} \kappa_k \Delta s_k, \\ (x, y)_{k+1} &= (x, y)_k + \Delta s_k (\cos \bar{\psi}_k, \sin \bar{\psi}_k), \\ \psi_{k+1} &= \psi_k + \kappa_k \Delta s_k, \end{aligned} \quad (1)$$

with $a_{\max} = 8 \text{ m/s}^2$, $\kappa_{\max} = 0.2 \text{ m}^{-1}$ and $a_{\text{lat}} = 6 \text{ m/s}^2$. The initial position is the last observed one. Speed and heading come from the last observed displacement, or the recorded heading is used below 0.5 m/s. Every mode is therefore drivable, and speed cannot become negative. The same head serves every focal type. The bounds cost little expressiveness: fitting controls directly to 2,000 vehicle-like training futures leaves a mean endpoint error of 0.16 m, with 0.55% of futures more than 2 m away.

TABLE I
VALIDATION RESULTS, ALL 24,988 FOCAL AGENTS. $K=1$ SCORES THE MOST PROBABLE MODE; “BRIER” IS BRIER-MINFDE, THE LEADERBOARD METRIC.

Model	$K=1$			$K=6$			
	minADE	minFDE	MR	minADE	minFDE	MR	brier
Constant velocity	4.40	11.66	.855	–	–	–	–
LSTM [18]	3.15	8.33	.821	–	–	–	–
Coordinate head	2.45	6.09	.709	1.06	1.94	.303	2.57
Kinematic head	2.74	6.95	.755	1.25	2.55	.431	3.16

TABLE II
PHYSICAL PLAUSIBILITY ON VALIDATION, VEHICLE-LIKE FOCAL AGENTS (23,113): SHARE OF TRAJECTORIES, MOST PROBABLE MODE / ALL SIX MODES.

	Infeasible, 10 Hz	Infeasible, 2 Hz	Off-road
Ground truth	2.93%	0.41%	0.19%
Coordinate head	57.8 / 70.0%	0.00 / 0.01%	1.52 / 1.75%
Kinematic head	0.00 / 0.00%	0.00 / 0.00%	5.17 / 5.68%

D. Training

Both heads use a winner-takes-all loss [13]. The mode with the lowest endpoint error gets the average displacement loss, and the logits get a cross-entropy loss towards that mode; the two losses are summed. Training uses Adam [17] with cosine decay, gradient clipping at 1.0 and batch size 64, for 30 epochs ($\approx 23.5\text{k}$ steps). The learning rates come from a three-point sweep on the 10% subset: 10^{-3} for the coordinate head, $3 \cdot 10^{-4}$ for the kinematic head. The checkpoint with the lowest dev brier-minFDE ($K=6$) is kept. A run takes 3.1–3.3 h on an Apple M4 laptop GPU, with the training subset held in memory. The data-efficiency runs keep the number of gradient steps fixed at $\approx 23.5\text{k}$ (300, 120 and 30 epochs) and check dev every ≈ 790 steps.

E. Metrics

Accuracy follows the Argoverse 2 definitions. Among the K most probable modes, the one with the lowest endpoint error is scored: minADE and minFDE are its mean and final errors, a miss (MR) is a minFDE above 2 m, and brier-minFDE adds $(1 - p)^2$ for that mode’s probability p . At $K=1$ the most probable mode is scored. Our implementation matches the `av2` package’s functions. Physical plausibility is measured for vehicle-like focal agents (vehicle, bus, motorcyclist). A trajectory is *infeasible* if its lateral acceleration (speed times yaw rate) exceeds 8 m/s^2 at any step while moving faster than 1 m/s. This is checked at 10 Hz and on the trajectory resampled to 2 Hz, which judges the manoeuvre rather than step-to-step noise. A trajectory is *off-road* if any point lies outside every drivable-area polygon of the scenario’s map. “@1” is the most probable mode and “@6” covers all modes.

TABLE III
ABLATIONS ON VALIDATION. BRIER-MINFDE AT THE MODEL’S OWN K (ALL FOCAL AGENTS; FOR $K=1$, MINFDE) AND OFF-ROAD FOR THE MOST PROBABLE MODE (VEHICLE-LIKE). GAPS ARE COMPUTED FROM UNROUNDED VALUES.

Setting	brier-minFDE (m)			Off-road @1	
	Coord.	Kinem.	Gap	Coord.	Kinem.
100%, $K=6$, seed 0	2.57	3.16	0.59	1.5%	5.2%
100%, $K=6$, seed 1	2.56	3.18	0.62	1.6%	5.9%
25% (12.5k scenes)	2.88	3.61	0.73	2.3%	7.9%
10% (5k scenes)	3.33	3.98	0.66	4.1%	8.1%
No map	3.18	3.58	0.41	4.0%	7.1%
$K=3$	3.17	3.72	0.54	1.6%	4.7%
$K=1$	5.19	5.81	0.62	2.0%	4.6%




Fig. 1. Data efficiency on validation. Every run gets $\approx 23.5k$ gradient steps, so smaller subsets get more epochs. The kinematic head’s gap does not shrink with less data.

IV. RESULTS

A. Accuracy and feasibility

Table I gives the main comparison. The Transformer with the coordinate head is 22% better than a focal-history LSTM [18] on $K=1$ minADE and 27% better on minFDE. The kinematic head is behind the coordinate head on every metric: brier-minFDE is 3.16 against 2.57, and $K=1$ minFDE is 6.95 against 6.09. It still beats the LSTM. Pedestrians do not explain the gap: for vehicle-like agents brier-minFDE is 3.28 against 2.64.

Table II answers the feasibility question. At 10 Hz, 58% of the coordinate head’s most probable trajectories break the lateral-acceleration limit. At that rate a ± 4 cm zigzag between steps already exceeds 8 m/s^2 , and even the tracked ground truth fails 2.9% of the time. Resampled to 2 Hz, the coordinate head is at 0.00%, against 0.41% for the ground truth: its manoeuvres are feasible, and its 10 Hz infeasibility is jitter. The kinematic head is 0.00% infeasible at both rates, as built. But it leaves the drivable area more than three times as often (5.2% against 1.5%).

B. Ablations

Table III varies one setting at a time, each a separate pair of training runs.

Seeds. A second seed moves each head’s brier-minFDE by at most 0.016 m, while the gap is 0.59 and 0.62 m.




Fig. 2. Mean $|\partial \mathcal{L} / \partial \text{head output}|$ per 1 s band of the future, relative to its mean over all 60 steps (dev, training loss). A control moves every later position, so integration makes the first second’s controls receive about 60 times the gradient of the last second’s.

Data efficiency. The gap is 0.66 m at 10%, 0.73 m at 25% and 0.59 m at 100%, or 20–25% of the coordinate head’s error (Fig. 1). The kinematic head trained on 50k scenes is behind the coordinate head trained on 12.5k. The one exception is the most probable mode at 10%, where the kinematic head is slightly ahead ($K=1$ minFDE 8.51 against 8.62). Both heads overfit the small subsets, the kinematic head earlier and further: at 10% its dev brier-minFDE rises from 4.02 at epoch 80 to 5.25 at epoch 300. Its training displacement error then falls to 0.73 m, the same as the coordinate head’s 0.74 m and far below the 1.18 m at which the full kinematic run ends. Given enough passes, it fits its training scenes as well as the coordinate head does, so the full-data gap reflects slow optimisation rather than missing expressiveness. The coordinate head’s jitter, meanwhile, grows as data shrinks: 97% of its most probable trajectories are infeasible at 10 Hz when trained on 10%.

Map. Removing the lane tokens costs the coordinate head 0.60 m and the kinematic head 0.42 m. The gap shrinks to 0.41 m but remains, so a weaker use of the map explains only part of it. With the map, the kinematic head still leaves the road more often than the coordinate head does without one (5.2% against 4.0%).

Number of modes. Models trained with $K=1$ and $K=3$ are selected on dev minFDE and brier-minFDE at $K=3$, respectively. The gap is 0.62, 0.54 and 0.59 m at $K=1$, 3 and 6. A single-mode predicts better than a six-mode model’s most probable mode (5.19 against 6.09 for the coordinate head, 5.81 against 6.95 for the kinematic head), so mode scoring is a weak point of both $K=6$ models.

V. DIAGNOSIS

The following analyses compare the two full-data checkpoints on the dev set.

A. Why the kinematic head is less accurate

Late controls get little gradient. Fig. 2 shows how the training loss’s gradient reaches the head’s outputs. For the coordinate head it is flat over the horizon. For the kinematic head it falls from 2.5 times the mean in the first second to 0.04 times in the last, for both acceleration and steering. This

TABLE IV

OFF-ROAD SHARE OF THE MOST PROBABLE MODE ON DEV, VEHICLE-LIKE AGENTS (4,636), BY CONDITION. HEADING CHANGE: FIRST AGAINST LAST 1 S CHORD OF THE GROUND TRUTH. e_0 : ANGLE BETWEEN THE KINEMATIC HEAD’S INITIAL HEADING AND THE GROUND TRUTH’S FIRST-SECOND DIRECTION. LATERAL MISS: ENDPOINT ERROR ACROSS THE GROUND TRUTH’S FINAL DIRECTION (SCENE COUNTS DIFFER PER HEAD).

Condition	Scenes	Coord.	Kinem.
All	4,636	1.6%	5.2%
Straight ($< 10^\circ$), $ e_0 < 1^\circ$	2,016	0.2%	1.2%
Straight ($< 10^\circ$), $ e_0 \geq 4^\circ$	125	2.4%	18.4%
Bend ($10\text{--}30^\circ$)	305	4.6%	14.8%
Turn ($\geq 30^\circ$)	584	6.8%	18.2%
Lateral miss 2–4 m	362 / 491	5.5%	10.6%
Lateral miss ≥ 4 m	295 / 516	9.2%	28.3%

is an effect of integrating through time, akin to the gradient problems of recurrent networks [19]. The best-of-six error gap grows with the horizon accordingly: 0.16 m at 3 s, 0.46 m at 5 s and 0.65 m at 6 s.

Accelerations chatter at their bounds. In the most probable mode, 31% of accelerations lie within 5% of $\pm 8 \text{ m/s}^2$ (55% in the last 2 s), against 6% of steering values. In the first 4 s the acceleration changes sign at 57–80% of steps. Over 1 s windows, though, it averages out close to the ground truth (median $|a|$ 0.88 against 0.61 m/s^2). Where tanh saturates, little gradient passes. Forcing smooth controls does not help: knots every 0.5 s with linear interpolation remove the chatter by construction and raise the fitting floor only from 0.16 to 0.17 m. On the 10% subset (30 epochs), however, dev brier-minFDE is worse with knots at both seeds (4.41 and 4.21 against 4.28 and 4.13). The imbalance comes from integration, not from the chatter.

The ground truth ends with an artifact. Over the last 0.7 s, the focal agent’s ground-truth positions slow down sharply while its recorded velocity does not. Vehicle-like speed from position differences falls from 7.7 to 3.6 m/s, an apparent -8.6 m/s^2 at the last step. The artifact is present in the raw data. A coordinate head can copy it, while the kinematic head would have to brake at its bound. It accounts for only about 0.1 m of the gap, however: the gap grows at the same rate before and after it begins.

B. Why the kinematic head leaves the road

Table IV and Fig. 3 locate the off-road excess.

Drivable is not road-aware. The bicycle model limits how the vehicle moves, not where it may go: a smooth, drivable path can still cross a median. Staying on the road is learned from the map, by both heads alike. What the constraint changes is how errors build up. The coordinate head places each point directly and can follow the lane’s shape point by point. The kinematic head’s position is the running sum of its controls, so a small early error in heading or steering compounds into a smooth drift across the lane.

Not the chatter. Re-rolling the kinematic head’s own controls after a 1 s moving average gives 5.0% off-road instead of 5.2%. Removing its steering in the last second gives 5.5%.

Mostly turns. Futures that change heading by 10° or more are 19% of the scenes but 62% of the kinematic head’s off-road modes. Its most probable mode turns less than the ground truth (median 0.85 of the heading change, against 0.92 for the coordinate head), and 61% of its off-road modes in turns end on the outside of the turn (coordinate head 55%). The bounds are not what stops it: its best mode turns as far as the coordinate head’s (0.92 against 0.94), and in the turns where it leaves the road its steering is at the bound in only 18% of steps.

Uncorrected initial heading. On straight roads, over its first second, the kinematic head’s most probable mode keeps 94% of the error e_0 in its initial heading, where the coordinate head keeps 3%. The last-step displacement is the better of the two available starting headings: at every moving speed it is closer to the first ground-truth step than the recorded heading is. Still, the head does not correct what error remains. Its off-road share on straight roads rises from 1.2% for $|e_0| < 1^\circ$ to 18.4% for $|e_0| \geq 4^\circ$ (coordinate head: 0.2% to 2.4%).

Errors across the lane. When the most probable mode misses the ground truth’s endpoint by more than 2 m sideways, the kinematic head leaves the road 2–3 times as often as the coordinate head with the same miss.

VI. DISCUSSION AND LIMITATIONS

For this encoder and benchmark, hard per-step kinematic decoding has a clear cost and a small benefit. The cost is 0.4–0.7 m of brier-minFDE in every setting, with no gain in data efficiency. The benefit is the removal of 10 Hz jitter; the coordinate head’s manoeuvres are already feasible at 2 Hz. The kinematic head can fit futures within 0.16 m and, given enough passes, fits its training scenes as well as the coordinate head. Its problems are therefore optimisation and initial state, not capacity. These point to remedies we have not tested: rescaling the gradient per step, predicting controls as residuals around a lane-following reference, estimating the initial state from more of the history, or projecting a coordinate decoder’s output onto feasible trajectories instead of decoding controls directly.

The study has clear limits. It uses one small model (1.43M parameters), a 50k-scene training subset, one benchmark and two seeds for the main comparison. Results are on validation only, because the leaderboard had closed. Learning rates come from a three-point sweep on the 10% subset. The kinematic head also serves pedestrians, although the gap is the same for vehicle-like agents. The diagnoses use one checkpoint per head. Off-road depends on the map’s drivable-area polygons.

VII. CONCLUSION

This project set out to build a physically accurate forecaster and measure how it performs, not to find the best approach. Integrating bounded controls through a kinematic bicycle model makes every trajectory drivable, but on Argoverse 2 it costs 0.4–0.7 m in every setting and does not improve data




Fig. 3. Dev scenes, each drawn at random from the scenes that meet its condition. Thick lines are the most probable modes. (a) A turn in which both heads’ most probable modes end within 2 m (one of 3). (b) A turn in which the kinematic head’s most probable mode turns the right way but only 25–75% as far, and leaves the road while the coordinate head’s stays on it (one of 19). (c) A straight road on which the kinematic head starts 4.1° off the ground truth’s direction and keeps that heading across the median (one of 21). Panels (b) and (c) show failure modes, not how often they occur: overall, the kinematic head’s most probable mode stays on the road in about 95% of vehicle-like scenes (5.2% off-road on both dev and validation).

efficiency. Correcting the initial heading and the slow learning of late controls, rather than adding constraints, is the natural next step.

REFERENCES

- [1] J. Gao, C. Sun, H. Zhao, Y. Shen, D. Anguelov, C. Li, and C. Schmid, “VectorNet: Encoding HD maps and agent dynamics from vectorized representation,” in *Proc. IEEE/CVF Conf. Computer Vision and Pattern Recognition (CVPR)*, 2020.
- [2] J. Ngiam, B. Caine, V. Vasudevan, Z. Zhang, H.-T. L. Chiang *et al.*, “Scene transformer: A unified architecture for predicting future trajectories of multiple agents,” in *Proc. Int. Conf. Learning Representations (ICLR)*, 2022.
- [3] N. Nayakanti, R. Al-Rfou, A. Zhou, K. Goel, K. S. Refaat, and B. Sapp, “Wayformer: Motion forecasting via simple & efficient attention networks,” in *Proc. IEEE Int. Conf. Robotics and Automation (ICRA)*, 2023.
- [4] S. Shi, L. Jiang, D. Dai, and B. Schiele, “Motion transformer with global intention localization and local movement refinement,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.
- [5] Z. Zhou, J. Wang, Y.-H. Li, and Y.-K. Huang, “Query-centric trajectory prediction,” in *Proc. IEEE/CVF Conf. Computer Vision and Pattern Recognition (CVPR)*, 2023.
- [6] H. Cui, T. Nguyen, F.-C. Chou, T.-H. Lin, J. Schneider, D. Bradley, and N. Djuric, “Deep kinematic models for kinematically feasible vehicle trajectory predictions,” in *Proc. IEEE Int. Conf. Robotics and Automation (ICRA)*, 2020.
- [7] T. Salzmann, B. Ivanovic, P. Chakravarty, and M. Pavone, “Trajectron++: Dynamically-feasible trajectory forecasting with heterogeneous data,” in *Proc. European Conf. Computer Vision (ECCV)*, 2020.
- [8] B. Wilson, W. Qi, T. Agarwal, J. Lambert, J. Singh *et al.*, “Argoverse 2: Next generation datasets for self-driving perception and forecasting,” in *Proc. NeurIPS Datasets and Benchmarks Track*, 2021.
- [9] J. Kong, M. Pfeiffer, G. Schildbach, and F. Borrelli, “Kinematic and dynamic vehicle models for autonomous driving control design,” in *Proc. IEEE Intelligent Vehicles Symposium (IV)*, 2015.
- [10] P. Polack, F. Althé, B. d’Andréa Novel, and A. de La Fortelle, “The kinematic bicycle model: A consistent model for planning feasible trajectories for autonomous vehicles?” in *Proc. IEEE Intelligent Vehicles Symposium (IV)*, 2017.
- [11] M. Liang, B. Yang, R. Hu, Y. Chen, R. Liao, S. Feng, and R. Urtasun, “Learning lane graph representations for motion forecasting,” in *Proc. European Conf. Computer Vision (ECCV)*, 2020.
- [12] C. R. Qi, H. Su, K. Mo, and L. J. Guibas, “PointNet: Deep learning on point sets for 3D classification and segmentation,” in *Proc. IEEE Conf. Computer Vision and Pattern Recognition (CVPR)*, 2017.
- [13] A. Guzman-Rivera, D. Batra, and P. Kohli, “Multiple choice learning: Learning to produce multiple structured outputs,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2012.
- [14] Y. Chai, B. Sapp, M. Bansal, and D. Anguelov, “MultiPath: Multiple probabilistic anchor trajectory hypotheses for behavior prediction,” in *Proc. Conf. Robot Learning (CoRL)*, 2019.
- [15] T. Phan-Minh, E. C. Grigore, F. A. Boulton, O. Beijbom, and E. M. Wolff, “CoverNet: Multimodal behavior prediction using trajectory sets,” in *Proc. IEEE/CVF Conf. Computer Vision and Pattern Recognition (CVPR)*, 2020.
- [16] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and I. Polosukhin, “Attention is all you need,” in *Advances in Neural Information Processing Systems (NeurIPS)*, 2017.
- [17] D. P. Kingma and J. Ba, “Adam: A method for stochastic optimization,” in *Proc. Int. Conf. Learning Representations (ICLR)*, 2015.
- [18] S. Hochreiter and J. Schmidhuber, “Long short-term memory,” *Neural Computation*, vol. 9, no. 8, pp. 1735–1780, 1997.
- [19] R. Pascanu, T. Mikolov, and Y. Bengio, “On the difficulty of training recurrent neural networks,” in *Proc. Int. Conf. Machine Learning (ICML)*, 2013.